

Pima Indians Diabetes - EDA & Prediction (0.776%)

- Accuracy - Decision Tree : 70.77%

Qasim Hassan

02-132182-028

Course: Artificial Intelligence

Submitted to : Ma'am Nabiha Faisal

July 2021

TABLE OF CONTENT

- [1. Load libraries and read the data](#)
 - [1.1. Load libraries](#)
 - [1.2. Read the data](#)
- [2. Overview](#)
 - [2.1. Head](#)
 - [2.2. Target](#)
 - [2.3. Missing values](#)
- [3. Replace missing values and EDA](#)
 - [3.1. Insulin](#)
 - [3.2. Glucose](#)
 - [3.3. SkinThickness](#)
 - [3.4. BloodPressure](#)
 - [3.5. BMI](#)
- [4. New features \(16\) and EDA](#)
- [5. Prepare dataset](#)
 - [5.1. StandardScaler and LabelEncoder](#)
 - [5.2. Correlation Matrix](#)
- [6. Machine Learning](#)
 - [6.1. Decision Tree Algorithm](#)
 - [6.2. How does the Decision Tree algorithm work?](#)
 - [6.3. Decision Tree Classifier Building in Scikit-learn](#)
 - [6.4. Model Performance](#)
- [7. Credits](#)

Who is Pima Indians ?

"The Pima (or Akimel O'odham, also spelled Akimel O'otham, "River People", formerly known as Pima) are a group of Native Americans living in an area consisting of what is now central and southern Arizona. The majority population of the surviving two bands of the Akimel O'odham are based in two reservations: the Kelli Akimel O'otham on the Gila River Indian Community (GRIC) and the On'k Akimel O'odham on the Salt River Pima-Maricopa Indian Community (SRPMIC)." Wikipedia

1. Load libraries and read the data

1.1. Load libraries

Loading the libraries

```
In [4]: 1 # Python Libraries
        2 # Classic, data manipulation and linear algebra
        3 import pandas as pd
        4 import numpy as np
        5
        6 # Plots
        7 import seaborn as sns
        8 import matplotlib.pyplot as plt
        9 %matplotlib inline
       10 import plotly.offline as py
       11 import plotly.graph_objs as go
       12 from plotly.offline import download_plotlyjs, init_notebook_mode, plot, iplot
       13 import plotly.tools as tls
       14 import plotly.figure_factory as ff
       15 py.init_notebook_mode(connected=True)
       16 import squarify
       17
```

```

18 # Data processing, metrics and modeling
19 from sklearn.preprocessing import StandardScaler, LabelEncoder
20 from sklearn.model_selection import GridSearchCV, cross_val_score, train_test_split, GridSearchCV, RandomizedSearchCV
21 from sklearn.metrics import precision_score, recall_score, confusion_matrix, roc_curve, precision_recall_curve, accuracy_score
22 import lightgbm as lgbm
23
24
25 # Stats
26 import scipy.stats as ss
27 from scipy import interp
28 from scipy.stats import randint as sp_randint
29 from scipy.stats import uniform as sp_uniform
30
31 # Time
32 from contextlib import contextmanager
33 @contextmanager
34 def timer(title):
35     t0 = time.time()
36     yield
37     print("{} - done in {:.0f}s".format(title, time.time() - t0))
38
39 #ignore warning messages
40 import warnings
41 warnings.filterwarnings('ignore')

```

1.2. Read data

Loading dataset with pandas (pd)

```
In [5]: 1 data = pd.read_csv('./data/')
```

2. Overview

2.1. Head

Checking data head and info

```
In [6]: 1 display(data.info(),data.head())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Pregnancies            768 non-null   int64
1   Glucose                 768 non-null   int64
2   BloodPressure           768 non-null   int64
3   SkinThickness           768 non-null   int64
4   Insulin                 768 non-null   int64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   int64
8   Outcome                 768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

None

```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

The datasets consist of several medical predictor (independent) variables and one target (dependent) variable, Outcome. Independent variables include the number of pregnancies the patient has had, their BMI, insulin level, age, and so on.

What is diabetes ?

According to NIH, "Diabetes is a disease that occurs when your blood glucose, also called blood sugar, is too high. Blood glucose is your main source of energy and comes from the food you eat. Insulin, a hormone made by the pancreas, helps glucose from food get into your cells to be used for energy. Sometimes your body doesn't make enough—or any—insulin or doesn't use insulin well. Glucose then stays in your blood and doesn't reach your cells.

Over time, having too much glucose in your blood can cause health problems. Although diabetes has no cure, you can take steps to manage your diabetes and stay healthy.

Sometimes people call diabetes "a touch of sugar" or "borderline diabetes." These terms suggest that someone doesn't really have diabetes or has a less serious case, but every case of diabetes is serious.

What are the different types of diabetes? The most common types of diabetes are type 1, type 2, and gestational diabetes.

Type 1 diabetes If you have type 1 diabetes, your body does not make insulin. Your immune system attacks and destroys the cells in your pancreas that make insulin. Type 1 diabetes is usually diagnosed in children and young adults, although it can appear at any age. People with type 1 diabetes need to take insulin every day to stay alive.

Type 2 diabetes If you have type 2 diabetes, your body does not make or use insulin well. You can develop type 2 diabetes at any age, even during childhood. However, this type of diabetes occurs most often in middle-aged and older people. Type 2 is the most common type of diabetes.

Gestational diabetes Gestational diabetes develops in some women when they are pregnant. Most of the time, this type of diabetes goes away after the baby is born. However, if you've had gestational diabetes, you have a greater chance of developing type 2 diabetes later in life. Sometimes diabetes diagnosed during pregnancy is actually type 2 diabetes.

Other types of diabetes Less common types include monogenic diabetes, which is an inherited form of diabetes, and cystic fibrosis-related diabetes ."

2.2 Target

What's target's distribution ?


The above graph shows that the data is unbalanced. The number of non-diabetic is 268 the number of diabetic patients is 500

```
In [7]: 1 # 2 datasets
2 D = data[(data['Outcome'] != 0)]
3 H = data[(data['Outcome'] == 0)]
4
5 #-----COUNT-----
6 def target_count():
7     trace = go.Bar( x = data['Outcome'].value_counts().values.tolist(),
8                     y = ['healthy', 'diabetic'],
9                     orientation = 'h',
10                    text=data['Outcome'].value_counts().values.tolist(),
11                    textfont=dict(size=15),
12                    textposition = 'auto',
13                    opacity = 0.8, marker=dict(
14                        color=['lightskyblue', 'gold'],
15                        line=dict(color='#000000', width=1.5)))
16
17     layout = dict(title = 'Count of Outcome variable')
18
19     fig = dict(data = [trace], layout=layout)
20     py.iplot(fig)
21
22 #-----PERCENTAGE-----
23 def target_percent():
24     trace = go.Pie(labels = ['healthy', 'diabetic'], values = data['Outcome'].value_counts(),
25                    textfont=dict(size=15), opacity = 0.8,
26                    marker=dict(colors=['lightskyblue', 'gold'],
27                                line=dict(color='#000000', width=1.5)))
28
29     layout = dict(title = 'Distribution of Outcome variable')
30
31     fig = dict(data = [trace], layout=layout)
32     py.iplot(fig)
33
```


```
In [8]: 1 target_count()
2 target_percent()
```



Count of Outcome variable



Distribution of Outcome variable



2.3. Missing values

We saw on `data.head()` that some features contain 0, it doesn't make sense here and this indicates missing value Below we replace 0 value by NaN :


```
In [9]: 1 data[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']] = data[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI'].replace(0, np.NaN)]
```

Now, we can look at where are missing values :

```
In [10]: 1 # Define missing plot to detect all missing values in dataset
2 def missing_plot(dataset, key) :
3     null_feat = pd.DataFrame(len(dataset[key]) - dataset.isnull().sum(), columns = ['Count'])
4     percentage_null = pd.DataFrame((len(dataset[key]) - (len(dataset[key]) - dataset.isnull().sum()))/len(dataset[key])*100,
5     percentage_null = percentage_null.round(2)
6
7     trace = go.Bar(x = null_feat.index, y = null_feat['Count'], opacity = 0.8, text = percentage_null['Count'], textposition = 'top',
8     line=dict(color='#000000',width=1.5))
9
10    layout = dict(title = "Missing Values (count & %)")
11
12    fig = dict(data = [trace], layout=layout)
13    py.iplot(fig)
14
```

```
In [11]: 1 # Plotting
2 missing_plot(data, 'Outcome')
```


Missing Values (count & %)



Missing values :

- Insulin = 48.7% - 374
- SkinThickness = 29.56% - 227
- BloodPressure = 4.56% - 35
- BMI = 1.43% - 11
- Glucose = 0.65% - 5

```
In [12]: 1 plt.style.use('ggplot') # Using ggplot2 style visuals
2
3 f, ax = plt.subplots(figsize=(11, 15))
4
5 ax.set_facecolor('#fafafa')
6 ax.set(xlim=(-.05, 200))
7 plt.ylabel('Variables')
8 plt.title("Overview Data Set")
9 ax = sns.boxplot(data = data,
10                 orient = 'h',
11                 palette = 'Set2')
```




OK, all missing values are encoded with NaN value

To fill these Nan values the data distribution needs to be understood against the target.

```
In [13]: 1 def correlation_plot():
2     #correlation
3     correlation = data.corr()
4     #tick labels
5     matrix_cols = correlation.columns.tolist()
6     #convert to array
7     corr_array = np.array(correlation)
8     trace = go.Heatmap(z = corr_array,
9                        x = matrix_cols,
10                       y = matrix_cols,
11                       colorscale='Viridis',
12                       colorbar = dict(),
13                       )
14     layout = go.Layout(dict(title = 'Correlation Matrix for variables',
15                             #autosize = False,
16                             #height = 1400,
17                             #width = 1600,
18                             margin = dict(r = 0, l = 100,
19                                           t = 0, b = 100,
20                                           ),
21                             yaxis = dict(tickfont = dict(size = 9)),
22                             xaxis = dict(tickfont = dict(size = 9)),
23                             )
24     fig = go.Figure(data = [trace], layout = layout)
25     py.iplot(fig)
```

A correlation matrix is a table showing correlation coefficients between sets of variables. Each random variable (X_i) in the table is correlated with each of the other values in the table (X_j). This allows you to see which pairs have the highest correlation.

```
In [14]: 1 correlation_plot()
```



To replace missing values, we'll use median by target (Outcome)

3. Replace missing values and EDA

```
In [15]: 1 def median_target(var):
2     temp = data[data[var].notnull()]
3     temp = temp[[var, 'Outcome']].groupby(['Outcome'])[var].median().reset_index()
4     return temp
```

3.1. Insulin

- Insulin : 2-Hour serum insulin (mu U/ml)

```
In [16]: 1 def plot_distribution(data_select, size_bin):
2     # 2 datasets
3     tmp1 = D[data_select]
4     tmp2 = H[data_select]
5     hist_data = [tmp1, tmp2]
6
7     group_labels = ['diabetic', 'healthy']
8     colors = ['#FD7000', '#7EC0EE']
9
10    fig = ff.create_distplot(hist_data, group_labels, colors = colors, show_hist = True, bin_size = size_bin, curve_type='kd
11
```


```

11 fig['layout'].update(title = data_select)
12
13
14 py.iplot(fig, filename = 'Density plot')

```

```
In [17]: 1 plot_distribution('Insulin', 0)
```

Insulin



```
In [18]: 1 median_target('Insulin')
```

Out[18]:

	Outcome	Insulin
0	0	102.5
1	1	169.5

Insulin's medians by the target are really different ! 102.5 for a healthy person and 169.5 for a diabetic person

```

In [19]: 1 data.loc[(data['Outcome'] == 0) & (data['Insulin'].isnull()), 'Insulin'] = 102.5
2 data.loc[(data['Outcome'] == 1) & (data['Insulin'].isnull()), 'Insulin'] = 169.5

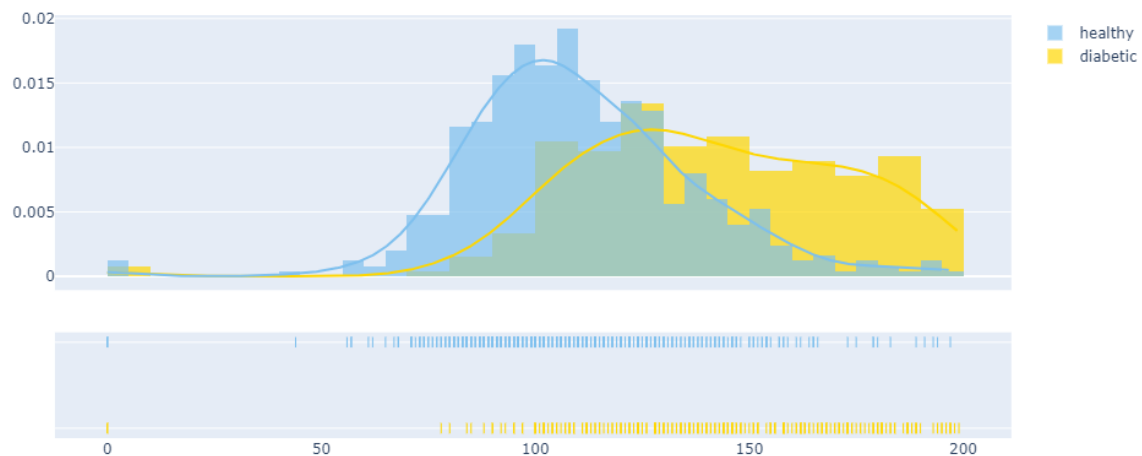
```

3.2. Glucose

- Glucose : Plasma glucose concentration a 2 hours in an oral glucose tolerance test

```
In [20]: 1 plot_distribution('Glucose', 0)
```

Glucose



```
In [21]: 1 median_target('Glucose')
```

```
Out[21]:
```

	Outcome	Glucose
0	0	107.0
1	1	140.0

```
In [22]: 1 data.loc[(data['Outcome'] == 0) & (data['Glucose'].isnull()), 'Glucose'] = 107
2 data.loc[(data['Outcome'] == 1) & (data['Glucose'].isnull()), 'Glucose'] = 140
```

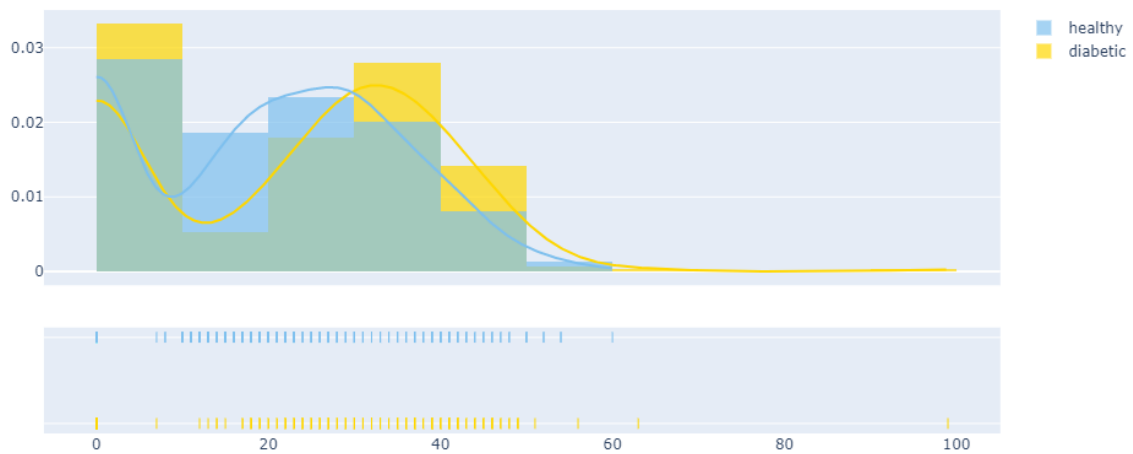
107 for a healthy person and 140 for a diabetic person

3.3. SkinThickness

- ****SkinThickness****: Triceps skin fold thickness (mm)

```
In [23]: 1 plot_distribution('SkinThickness', 10)
```

SkinThickness



```
In [24]: 1 median_target('SkinThickness')
```

```
Out[24]:
```

	Outcome	SkinThickness
0	0	27.0
1	1	32.0

```
In [25]: 1 data.loc[(data['Outcome'] == 0) & (data['SkinThickness'].isnull()), 'SkinThickness'] = 27
2 data.loc[(data['Outcome'] == 1) & (data['SkinThickness'].isnull()), 'SkinThickness'] = 32
```


27 for a healthy person and 32 for a diabetic person

3.4. BloodPressure

- ****BloodPressure****: Diastolic blood pressure (mm Hg)

```
In [26]: 1 plot_distribution('BloodPressure', 5)
```


BloodPressure



```
In [27]: 1 median_target('BloodPressure')
```

```
Out[27]:
```

	Outcome	BloodPressure
0	0	70.0
1	1	74.5


```
In [28]: 1 data.loc[(data['Outcome'] == 0) & (data['BloodPressure'].isnull()), 'BloodPressure'] = 70
         2 data.loc[(data['Outcome'] == 1) & (data['BloodPressure'].isnull()), 'BloodPressure'] = 74.5
```

3.5. BMI

- BMI : Body mass index (weight in kg/(height in m)²)

```
In [29]: 1 plot_distribution('BMI', 0)
```

BMI



```
In [30]: 1 median_target('BMI')
```

```
Out[30]:
```


	Outcome	BMI
0	0	30.1
1	1	34.3

```
In [31]: 1 data.loc[(data['Outcome'] == 0) & (data['BMI'].isnull()), 'BMI'] = 30.1
2 data.loc[(data['Outcome'] == 1) & (data['BMI'].isnull()), 'BMI'] = 34.3
```


- **Age** : Age (years)
- **DiabetesPedigreeFunction** : Diabetes pedigree function
- **Pregnancies** : Number of times pregnant

```
In [32]: 1 #plot distribution
2 plot_distribution('Age', 0)
3 plot_distribution('Pregnancies', 0)
4 plot_distribution('DiabetesPedigreeFunction', 0)
```


Age



Pregnancies




DiabetesPedigreeFunction



```
In [33]: 1 missing_plot(data, 'Outcome')
```

Missing Values (count & %)



4. New features (16) and EDA

Here, we define 3 plots functions

```
In [34]: 1 def plot_feat1_feat2(featt1, featt2) :
2         D = data[(data['Outcome'] != 0)]
3         H = data[(data['Outcome'] == 0)]
4         trace0 = go.Scatter(
5             x = D[featt1],
6             y = D[featt2],
7             name = 'diabetic',
8             mode = 'markers',
9             marker = dict(color = '#FFD700',
10                line = dict(
11                    width = 1)))
12
13         trace1 = go.Scatter(
14             x = H[featt1],
15             y = H[featt2],
16             name = 'healthy',
17             mode = 'markers',
18             marker = dict(color = '#7EC0EE',
19                line = dict(
20                    width = 1)))
21
22         layout = dict(title = featt1 + " "+"vs"+" " + featt2,
23             yaxis = dict(title = featt2,zeroline = False),
24             xaxis = dict(title = featt1, zeroline = False)
25             )
26
27         plots = [trace0, trace1]
28
29         fig = dict(data = plots, layout=layout)
30         py.iplot(fig)

In [35]: 1 def barplot(var_select, sub) :
2         tmp1 = data[(data['Outcome'] != 0)]
3         tmp2 = data[(data['Outcome'] == 0)]
4         tmp3 = pd.DataFrame(pd.crosstab(data[var_select],data['Outcome']), )
5         tmp3['% diabetic'] = tmp3[1] / (tmp3[1] + tmp3[0]) * 100
6
7         color=['lightskyblue','gold' ]
8         trace1 = go.Bar(
9             x=tmp1[var_select].value_counts().keys().tolist(),
10            y=tmp1[var_select].value_counts().values.tolist(),
11            text=tmp1[var_select].value_counts().values.tolist(),
12            textposition = 'auto',
13            name='% diabetic',opacity = 0.8, marker=dict(
14                color='gold',
15                line=dict(color='#000000',width=1)))
16
17         trace2 = go.Bar(
18             x=tmp2[var_select].value_counts().keys().tolist(),
19             y=tmp2[var_select].value_counts().values.tolist(),
20             text=tmp2[var_select].value_counts().values.tolist(),
21             textposition = 'auto',
22             name='healthy', opacity = 0.8, marker=dict(
23                 color='lightskyblue',
24                 line=dict(color='#000000',width=1)))
25
26         trace3 = go.Scatter(
27             x=tmp3.index,
28             y=tmp3['% diabetic'],
29             yaxis = 'y2',
30             name='% diabetic', opacity = 0.6, marker=dict(
31                 color='black',
32                 line=dict(color='#000000',width=0.5
33                 )))
34
35         layout = dict(title = str(var_select)+' '(sub),
36             xaxis=dict(),
37             yaxis=dict(title= 'Count'),
38             yaxis2=dict(range= [-0, 75],
39                 overlaying= 'y',
40                 anchor= 'x',
41                 side= 'right',
42                 zeroline=False,
43                 showgrid= False,
44                 title= '% diabetic'
45             ))
46
47         fig = go.Figure(data=[trace1, trace2, trace3], layout=layout)
48         py.iplot(fig)
```

```

In [36]: 1 # Define pie plot to visualize each variable repartition vs target modalities : Survived or Died (train)
2 def plot_pie(var_select, sub):
3     D = data[(data['Outcome'] != 0)]
4     H = data[(data['Outcome'] == 0)]
5
6     col = ['Silver', 'mediumturquoise', '#CF5C36', 'lightblue', 'magenta', '#FF5D73', '#F2D7EE', 'mediumturquoise']
7
8     trace1 = go.Pie(values = D[var_select].value_counts().values.tolist(),
9                    labels = D[var_select].value_counts().keys().tolist(),
10                   textfont=dict(size=15), opacity = 0.8,
11                   hole = 0.5,
12                   hoverinfo = "label+percent+name",
13                   domain = dict(x = [.0, .48]),
14                   name = "Diabetic",
15                   marker = dict(colors = col, line = dict(width = 1.5)))
16     trace2 = go.Pie(values = H[var_select].value_counts().values.tolist(),
17                    labels = H[var_select].value_counts().keys().tolist(),
18                   textfont=dict(size=15), opacity = 0.8,
19                   hole = 0.5,
20                   hoverinfo = "label+percent+name",
21                   marker = dict(line = dict(width = 1.5)),
22                   domain = dict(x = [.52, 1]),
23                   name = "Healthy")
24
25     layout = go.Layout(dict(title = var_select + " distribution by target <br>" + (sub),
26                            annotations = [ dict(text = "Diabetic" : "+"268",
27                                                  font = dict(size = 13),
28                                                  showarrow = False,
29                                                  x = .22, y = -0.1),
30                                              dict(text = "Healthy" : "+"500",
31                                                  font = dict(size = 13),
32                                                  showarrow = False,
33                                                  x = .8, y = -.1)]))
34
35     fig = go.Figure(data = [trace1, trace2], layout = layout)
36     py.iplot(fig)
37

```


- Glucose and Age

```

In [37]: 1 plot_feat1_feat2('Glucose', 'Age')

```

Glucose vs Age




Healthy persons are concentrate with an age <= 30 and glucose <= 120

```

In [38]: 1 palette = {0 : 'lightblue', 1 : 'gold'}
2 edgcolor = 'black'
3
4 fig = plt.figure(figsize=(12,8))
5
6 ax1 = sns.scatterplot(x = data['Glucose'], y = data['Age'], hue = "Outcome",
7                      data = data, palette = palette, edgcolor=edgcolor)
8
9 plt.annotate('N1', size=25, color='black', xy=(80, 30), xytext=(60, 35),
10            arrowprops=dict(facecolor='black', shrink=0.05),
11            )
12 plt.plot([50, 120], [30, 30], linewidth=2, color = 'red')
13 plt.plot([120, 120], [20, 30], linewidth=2, color = 'red')
14 plt.plot([50, 120], [20, 20], linewidth=2, color = 'red')
15 plt.plot([50, 50], [20, 30], linewidth=2, color = 'red')
16 plt.title('Glucose vs Age')
17 plt.show()


```



```
In [39]: 1 data.loc[:, 'N1']=0
          2 data.loc[(data['Age']<=30) & (data['Glucose']<=120), 'N1']=1
```


```
In [40]: 1 barplot('N1', ':Glucose <= 120 and Age <= 30')
```

N1 :Glucose <= 120 and Age <= 30



```
In [41]: 1 plot_pie('N1', '(Glucose <= 120 and Age <= 30)')
```

N1 distribution by target
(Glucose <= 120 and Age <= 30)



- **BMI**


According to wikipedia "The body mass index (BMI) or Quetelet index is a value derived from the mass (weight) and height of an individual. The BMI is defined as the body mass divided by the square of the body height, and is universally expressed in units of kg/m², resulting from mass in kilograms and height in metres."

30 kg/m² is the limit to obesity

```
In [42]: 1 data.loc[:, 'N2'] = 0
         2 data.loc[(data['BMI'] <= 30), 'N2'] = 1
```


```
In [43]: 1 barplot('N2', ': BMI <= 30')
```

N2 : BMI <= 30



```
In [44]: 1 plot_pie('N2', 'BMI <= 30')
```


N2 distribution by target
BMI <= 30



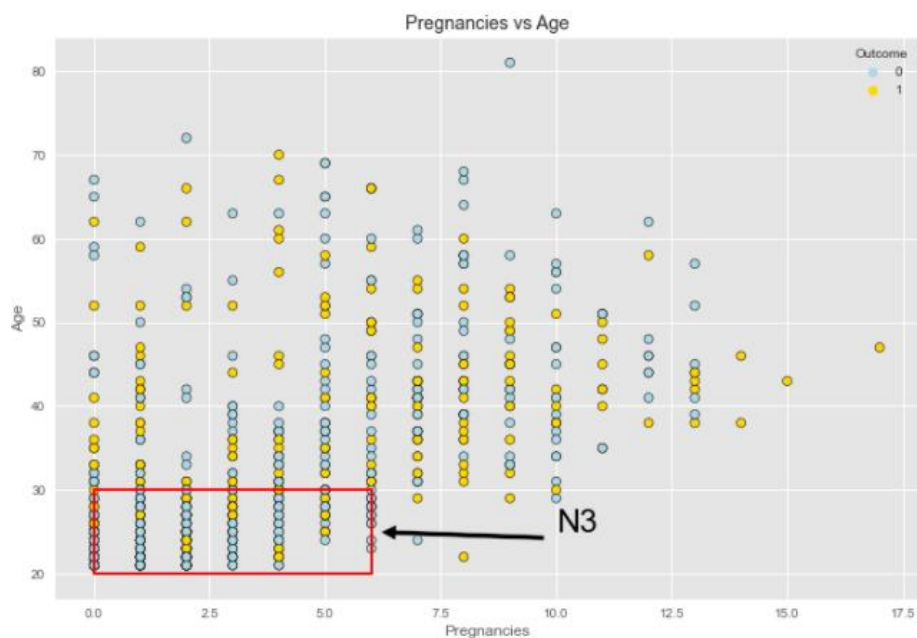
- **Pregnancies and Age**

```
In [45]: 1 plot_feat1_feat2('Pregnancies', 'Age')
```

Pregnancies vs Age




```
In [46]: 1 palette = {0 : 'lightblue', 1 : 'gold'}
2 edgecolor = 'black'
3
4 fig = plt.figure(figsize=(12,8))
5
6 ax1 = sns.scatterplot(x = data['Pregnancies'], y = data['Age'], hue = "Outcome",
7                       data = data, palette = palette, edgecolor=edgecolor)
8
9 plt.annotate('N3', size=25, color='black', xy=(6, 25), xytext=(10, 25),
10             arrowprops=dict(facecolor='black', shrink=0.05),
11             )
12 plt.plot([0, 6], [30, 30], linewidth=2, color = 'red')
13 plt.plot([6, 6], [20, 30], linewidth=2, color = 'red')
14 plt.plot([0, 6], [20, 20], linewidth=2, color = 'red')
15 plt.plot([0, 0], [20, 30], linewidth=2, color = 'red')
16 plt.title('Pregnancies vs Age')
17 plt.show()
```



```
In [47]: 1 data.loc[:, 'N3'] = 0
2 data.loc[(data['Age'] <= 30) & (data['Pregnancies'] <= 6), 'N3'] = 1
```


```
In [48]: 1 barplot('N3', ': Age <= 30 and Pregnancies <= 6')
```


N3 : Age <= 30 and Pregnancies <= 6



In [49]: 1 plot_pie('N3', 'Age <= 30 and Pregnancies <= 6')


N3 distribution by target
Age <= 30 and Pregnancies <= 6



• Glucose and BloodPressure

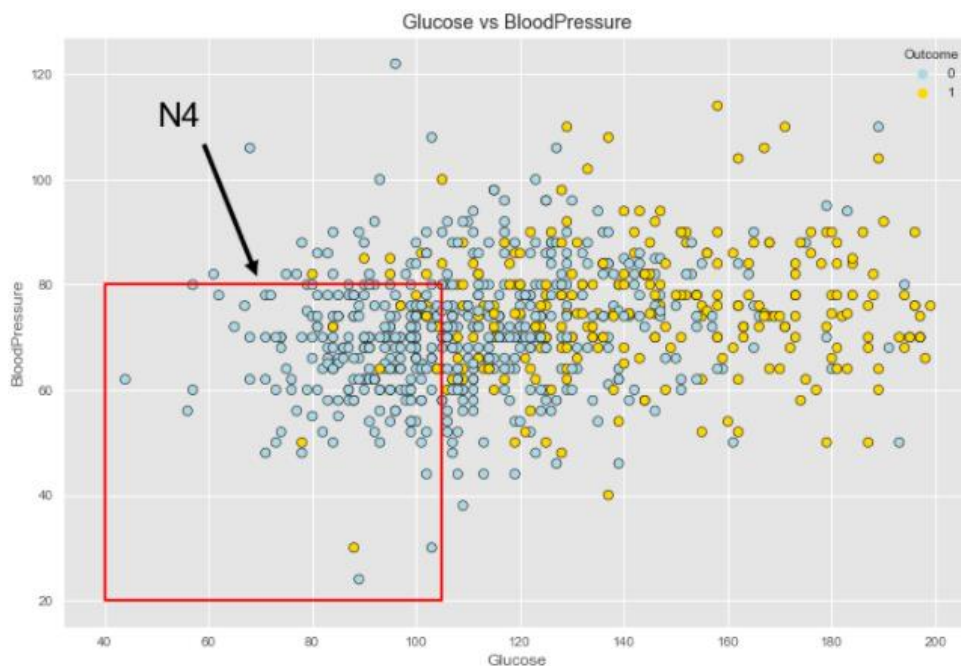
In [50]: 1 plot_feat1_feat2('Glucose', 'BloodPressure')

Glucose vs BloodPressure



Healthy persons are concentrate with an blood pressure ≤ 80 and glucose ≤ 105


```
In [51]: 1 palette = {0 : 'lightblue', 1 : 'gold'}
2 edgecolor = 'black'
3
4 fig = plt.figure(figsize=(12,8))
5
6 ax1 = sns.scatterplot(x = data['Glucose'], y = data['BloodPressure'], hue = "Outcome",
7                      data = data, palette = palette, edgecolor=edgecolor)
8
9 plt.annotate('N4', size=25, color='black', xy=(70, 80), xytext=(50, 110),
10            arrowprops=dict(facecolor='black', shrink=0.05),
11            )
12 plt.plot([40, 105], [80, 80], linewidth=2, color = 'red')
13 plt.plot([40, 40], [20, 80], linewidth=2, color = 'red')
14 plt.plot([40, 105], [20, 20], linewidth=2, color = 'red')
15 plt.plot([105, 105], [20, 80], linewidth=2, color = 'red')
16 plt.title('Glucose vs BloodPressure')
17 plt.show()
```



```
In [52]: 1 data.loc[:, 'N4'] = 0
2 data.loc[(data['Glucose'] <= 105) & (data['BloodPressure'] <= 80), 'N4'] = 1
```


```
In [53]: 1 barplot('N4', ': Glucose <= 105 and BloodPressure <= 80')
```

N4 : Glucose ≤ 105 and BloodPressure ≤ 80



```
In [54]: 1 plot_pie('N4', 'Glucose <= 105 and BloodPressure <= 80')
```

N4 distribution by target
Glucose ≤ 105 and BloodPressure ≤ 80




- SkinThickness

```
In [55]: 1 data.loc[:, 'N5'] = 0
          2 data.loc[(data['SkinThickness'] <= 20), 'N5'] = 1
```


```
In [56]: 1 barplot('N5', ':SkinThickness <= 20')
```

N5 :SkinThickness ≤ 20



```
In [57]: 1 plot_pie('N5', 'SkinThickness <= 20')
```

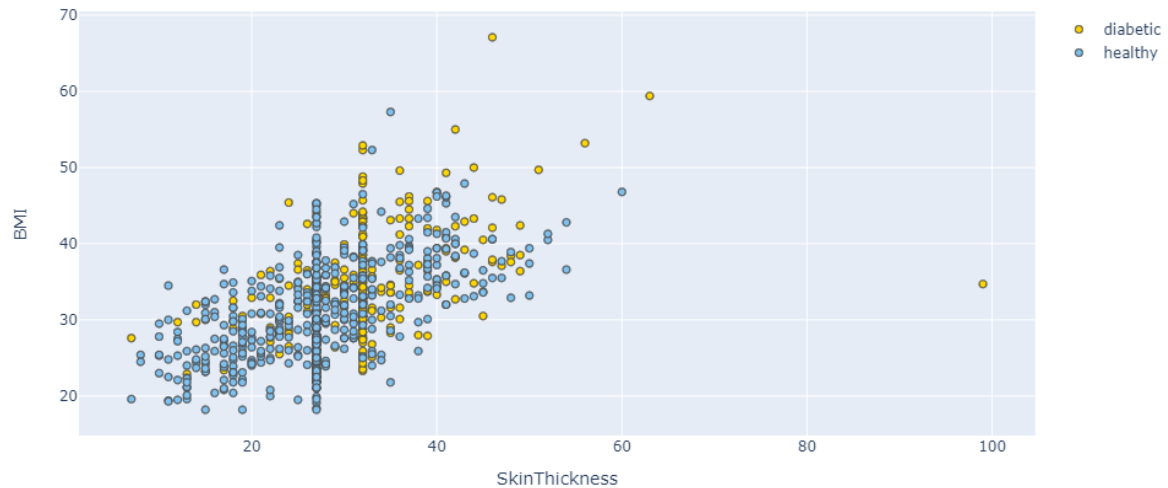
N5 distribution by target
SkinThickness ≤ 20



- SkinThickness and BMI

```
In [58]: 1 plot_feat1_feat2('SkinThickness', 'BMI')
```


SkinThickness vs BMI



Healthy persons are concentrate with a BMI < 30 and skin thickness <= 20


```
In [59]: 1 data.loc[:, 'N6'] = 0
2 data.loc[(data['BMI'] < 30) & (data['SkinThickness'] <= 20), 'N6'] = 1

In [60]: 1 palette = {0 : 'lightblue', 1 : 'gold'}
2 edgecolor = 'black'
3
4 fig = plt.figure(figsize=(12,8))
5
6 ax1 = sns.scatterplot(x = data['SkinThickness'], y = data['BMI'], hue = "Outcome",
7 data = data, palette = palette, edgecolor=edgecolor)
8
9 plt.annotate('N6', size=25, color='black', xy=(20, 20), xytext=(50, 25),
10 arrowprops=dict(facecolor='black', shrink=0.05),
11 )
12 plt.plot([0, 20], [30, 30], linewidth=2, color = 'red')
13 plt.plot([0, 0], [16, 30], linewidth=2, color = 'red')
14 plt.plot([0, 20], [16, 16], linewidth=2, color = 'red')
15 plt.plot([20, 20], [16, 30], linewidth=2, color = 'red')
16 plt.title('SkinThickness vs BMI')
17 plt.show()
```




```
In [61]: 1 barplot('N6', ': BMI < 30 and SkinThickness <= 20')
```

N6 : BMI < 30 and SkinThickness <= 20



```
In [62]: 1 plot_pie('N6', 'BMI < 30 and SkinThickness <= 20')
```


N6 distribution by target
BMI < 30 and SkinThickness <= 20




- Glucose and BMI

```
In [63]: 1 plot_feat1_feat2('Glucose', 'BMI')
```

Glucose vs BMI




```
In [64]: 1 palette = {0 : 'lightblue', 1 : 'gold'}
2         edgecolor = 'black'
3
4         fig = plt.figure(figsize=(12,8))
5
6         ax1 = sns.scatterplot(x = data['Glucose'], y = data['BMI'], hue = "Outcome",
7                             data = data, palette = palette, edgecolor=edgecolor)
8
9         plt.annotate('N7', size=25, color='black', xy=(70, 35), xytext=(40, 60),
10                    arrowprops=dict(facecolor='black', shrink=0.05),
11                    )
12         plt.plot([105, 105], [16, 30], linewidth=2, color = 'red')
13         plt.plot([40, 40], [16, 30], linewidth=2, color = 'red')
14         plt.plot([40, 105], [16, 16], linewidth=2, color = 'red')
15         plt.plot([40, 105], [30, 30], linewidth=2, color = 'red')
16         plt.title('Glucose vs BMI')
17         plt.show()
```



```
In [65]: 1 data.loc[:, 'N7'] = 0
2         data.loc[(data['Glucose'] <= 105) & (data['BMI'] <= 30), 'N7'] = 1
```


```
In [66]: 1 barplot('N7', ': Glucose <= 105 and BMI <= 30')
```

N7 : Glucose <= 105 and BMI <= 30



```
In [67]: 1 plot_pie('N7', 'Glucose <= 105 and BMI <= 30')
```


N7 distribution by target
Glucose ≤ 105 and BMI ≤ 30



• Insulin

```
In [68]: 1 plot_distribution('Insulin', 0)
```


Insulin



```
In [69]: 1 data.loc[:, 'N9'] = 0
         2 data.loc[(data['Insulin'] < 200), 'N9'] = 1
```


```
In [70]: 1 barplot('N9', ': Insulin < 200')
```

N9 : Insulin < 200



```
In [71]: 1 plot_pie('N9', 'Insulin < 200')
```

N9 distribution by target
Insulin < 200




• BloodPressure

```
In [72]: 1 data.loc[:, 'N10'] = 0
         2 data.loc[(data['BloodPressure'] < 80), 'N10'] = 1
```


```
In [73]: 1 barplot('N10', ': BloodPressure < 80')
```

N10 : BloodPressure < 80



```
In [74]: 1 plot_pie('N10', 'BloodPressure < 80')
```


N10 distribution by target
BloodPressure < 80



- Pregnancies

```
In [75]: 1 plot_distribution('Pregnancies', 0)
```


Pregnancies




```
In [76]: 1 data.loc[:, 'N11']=0
          2 data.loc[(data['Pregnancies']<4) & (data['Pregnancies']!=0), 'N11']=1
```

```
In [77]: 1 barplot('N11', ': Pregnancies > 0 and < 4')
```

N11 : Pregnancies > 0 and < 4



```
In [78]: 1 plot_pie('N11', 'Pregnancies > 0 and < 4')
```

N11 distribution by target
Pregnancies > 0 and < 4


- Others

```
In [79]: 1 data['N0'] = data['BMI'] * data['SkinThickness']
2
3 data['N8'] = data['Pregnancies'] / data['Age']
4
5 data['N13'] = data['Glucose'] / data['DiabetesPedigreeFunction']
6
7 data['N12'] = data['Age'] * data['DiabetesPedigreeFunction']
8
9 data['N14'] = data['Age'] / data['Insulin']
10
```

```
In [80]: 1 D = data[(data['Outcome'] != 0)]
2         H = data[(data['Outcome'] == 0)]
```

- ** Did you watch Inception ? ** Here is the same! It's not a dream in a dream but a new feature extract from a new feature


```
In [81]: 1 plot_distribution('N0', 0)
```



```
In [82]: 1 data.loc[:, 'N15'] = 0
2         data.loc[(data['N0'] < 1034), 'N15'] = 1
```


```
In [83]: 1 barplot('N15', ': N0 < 1034')
```

N15 : N0 < 1034



```
In [84]: 1 plot_pie('N15', 'N0 < 1034')
```

N15 distribution by target
N0 < 1034



5. Prepare dataset

5.1. StandardScaler and LabelEncoder

- **StandardScaler**:

Standardize features by removing the mean and scaling to unit variance :

Centering and scaling happen independently on each feature by computing the relevant statistics on the samples in the set. Mean and standard deviation are then stored to be used on later data using the transform method.

Standardization of a dataset is a common requirement for many machine learning estimators: they might behave badly if the individual features do not more or less look like standard normally distributed data (e.g. Gaussian with 0 mean and unit variance).

- **LabelEncoder**: Encode labels with value between 0 and n_classes-1.

Bellow we encode the data to feed properly to our algorithm

```
In [85]: 1 target_col = ["Outcome"]
2 cat_cols = data.nunique()[data.nunique() < 12].keys().tolist()
3 cat_cols = [x for x in cat_cols ]
4 #numerical columns
5 num_cols = [x for x in data.columns if x not in cat_cols + target_col]
6 #Binary columns with 2 values
7 bin_cols = data.nunique()[data.nunique() == 2].keys().tolist()
8 #Columns more than 2 values
9 multi_cols = [i for i in cat_cols if i not in bin_cols]
10
11 #Label encoding Binary columns
12 le = LabelEncoder()
13 for i in bin_cols :
14     data[i] = le.fit_transform(data[i])
15
16 #Duplicating columns for multi value columns
17 data = pd.get_dummies(data = data, columns = multi_cols )
18
19 #Scaling Numerical columns
20 std = StandardScaler()
21 scaled = std.fit_transform(data[num_cols])
22 scaled = pd.DataFrame(scaled, columns=num_cols)
23
24 #dropping original values merging scaled values for numerical columns
25 df_data Og = data.copy()
26 data = data.drop(columns = num_cols,axis = 1)
27 data = data.merge(scaled, left_index=True, right_index=True, how = "left")
```

Now, we can compute correlation matrix

5.2. Correlation Matrix

A **correlation matrix** is a table showing correlation coefficients between sets of variables. Each random variable (Xi) in the table is correlated with each of the other values in the table (Xj). This allows you to see which pairs have the highest correlation.


```
In [86]: 1 def correlation_plot():
2     #correlation
3     correlation = data.corr()
4     #tick labels
```

```

5 matrix_cols = correlation.columns.tolist()
6 #convert to array
7 corr_array = np.array(correlation)
8 trace = go.Heatmap(z = corr_array,
9                   x = matrix_cols,
10                  y = matrix_cols,
11                  colorscale='Viridis',
12                  colorbar = dict(),
13                  )
14 layout = go.Layout(dict(title = 'Correlation Matrix for variables',
15                        #autosize = False,
16                        #height = 1400,
17                        #width = 1600,
18                        margin = dict(r = 0, l = 100,
19                                t = 0, b = 100,
20                                ),
21                        yaxis = dict(tickfont = dict(size = 9)),
22                        xaxis = dict(tickfont = dict(size = 9)),
23                        )
24 )
25 fig = go.Figure(data = [trace], layout = layout)
26 py.ipplot(fig)
27

```

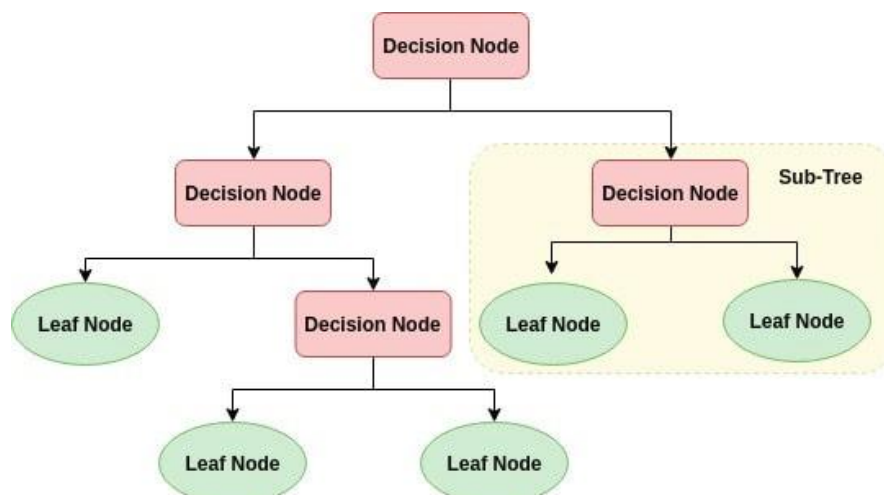
In [87]: 1 correlation_plot()



6. Machine Learning

6.1. Decision Tree

A decision tree is a flowchart-like tree structure where an internal node represents feature(or attribute), the branch represents a decision rule, and each leaf node represents the outcome. The topmost node in a decision tree is known as the root node. It learns to partition on the basis of the attribute value. It partitions the tree in a recursive manner called recursive partitioning. This flowchart-like structure helps you in decision making. It can be visualized like a flowchart diagram which easily mimics human level thinking. That is why decision trees are easy to understand and interpret.



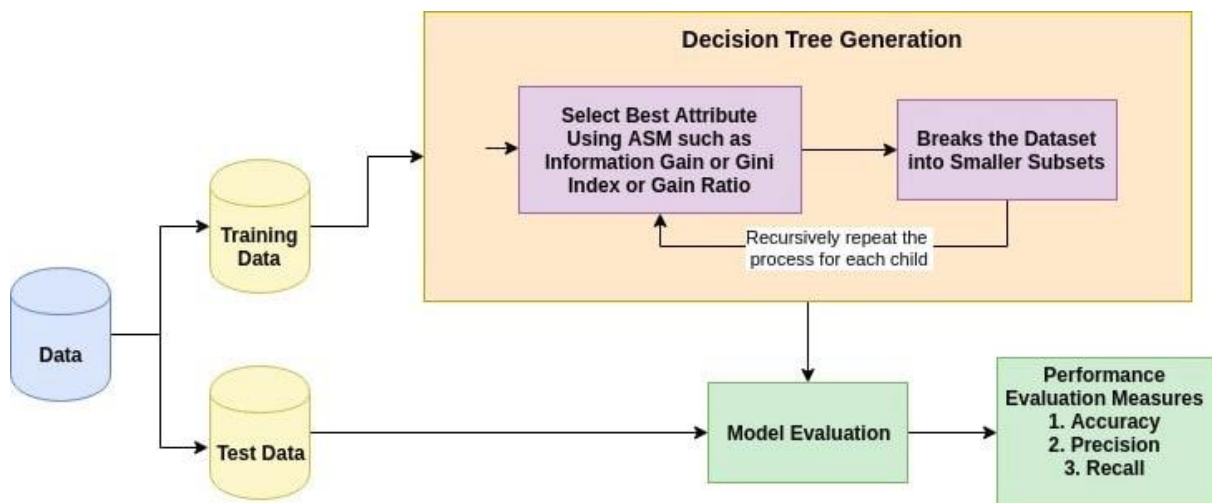
Decision Tree is a white box type of ML algorithm. It shares internal decision-making logic, which is not available in the black box type of algorithms such as Neural Network. Its training time is faster compared to the neural network algorithm. The time complexity of decision trees is a function of the number of records and number of attributes in the given data. The decision tree is a distribution-free or non-parametric method, which does not depend upon probability distribution assumptions. Decision trees can handle high dimensional data with good accuracy.

6.2. How does the Decision Tree algorithm work?

- **Discrimination Threshold** : A visualization of precision, recall, f1 score, and queue rate with respect to the discrimination threshold of a binary classifier. The discrimination threshold is the probability or score at which the positive class is chosen over the negative class

The basic idea behind any decision tree algorithm is as follows:

1. Select the best attribute using Attribute Selection Measures(ASM) to split the records.
2. Make that attribute a decision node and breaks the dataset into smaller subsets.
3. Start tree building by repeating this process recursively for each child until one of the condition will match:
 - All the tuples belong to the same attribute value.
 - There are no more remaining attributes.
 - There are no more instances.



Attribute Selection Measures

Attribute selection measure is a heuristic for selecting the splitting criterion that partition data into the best possible manner. It is also known as splitting rules because it helps us to determine breakpoints for tuples on a given node. ASM provides a rank to each feature(or attribute) by explaining the given dataset. Best score attribute will be selected as a splitting attribute ([Source](#)). In the case of a continuous-valued attribute, split points for branches also need to define. Most popular selection measures are Information Gain, Gain Ratio, and Gini Index.

Information Gain

Shannon invented the concept of entropy, which measures the impurity of the input set. In physics and mathematics, entropy referred as the randomness or the impurity in the system. In information theory, it refers to the impurity in a group of examples. Information gain is the decrease in entropy. Information gain computes the difference between entropy before split and average entropy after split of the dataset based on given attribute values. ID3 (Iterative Dichotomiser) decision tree algorithm uses information gain.

$$\text{Info}(D) = - \sum_{i=1}^m p_i \log_2 p_i$$

Where, P_i is the probability that an arbitrary tuple in D belongs to class C_i .

$$\text{Info}_A(D) = \sum_{j=1}^V \frac{|D_j|}{|D|} \times \text{Info}(D_j)$$

$$\text{Gain}(A) = \text{Info}(D) - \text{Info}_A(D)$$

Where:

- $\text{Info}(D)$ is the average amount of information needed to identify the class label of a tuple in D .
- $|D_j|/|D|$ acts as the weight of the j th partition.
- $\text{Info}_A(D)$ is the expected information required to classify a tuple from D based on the partitioning by A .

The attribute A with the highest information gain, $\text{Gain}(A)$, is chosen as the splitting attribute at node $N()$.

Gain Ratio

Information gain is biased for the attribute with many outcomes. It means it prefers the attribute with a large number of distinct values. For instance, consider an attribute with a unique identifier such as customer_ID has zero info(D) because of pure partition. This maximizes the information gain and creates useless partitioning.

C4.5, an improvement of ID3, uses an extension to information gain known as the gain ratio. Gain ratio handles the issue of bias by normalizing the information gain using Split Info. Java implementation of the C4.5 algorithm is known as J48, which is available in WEKA data mining tool.

$$\text{SplitInfo}_A(D) = - \sum_{j=1}^V \frac{|D_j|}{|D|} \times \log_2 \left(\frac{|D_j|}{|D|} \right)$$

Where:

- $|D_j|/|D|$ acts as the weight of the jth partition.
- v is the number of discrete values in attribute A.

The gain ratio can be defined as:

$$\text{GainRatio}(A) = \frac{\text{Gain}(A)}{\text{SplitInfo}_A(D)}$$

The attribute with the highest gain ratio is chosen as the splitting attribute ([Source](#)).

Gini index

Another decision tree algorithm CART (Classification and Regression Tree) uses the Gini method to create split points.

$$\text{Gini}(D) = 1 - \sum_{i=1}^m p_i^2$$

Where, p_i is the probability that a tuple in D belongs to class C_i .

The Gini Index considers a binary split for each attribute. You can compute a weighted sum of the impurity of each partition. If a binary split on attribute A partitions data D into D_1 and D_2 , the Gini index of D is:

$$\text{Gini}_A(D) = \frac{|D_1|}{D} \text{Gini}(D_1) + \frac{|D_2|}{D} \text{Gini}(D_2)$$

In case of a discrete-valued attribute, the subset that gives the minimum gini index for that chosen is selected as a splitting attribute. In the case of continuous-valued attributes, the strategy is to select each pair of adjacent values as a possible split-point and point with smaller gini index chosen as the splitting point.

$$\Delta \text{Gini}(A) = \text{Gini}(D) - \text{Gini}_A(D)$$

The attribute with minimum Gini index is chosen as the splitting attribute.

6.3. Decision Tree Classifier Building in Scikit-learn

Importing Required Libraries

Let's first install the required packages and load the required libraries.

```
In [92]: 1 #!pip install --quiet -r requirements.txt;

In [93]: 1 # Load Libraries
          2 import pandas as pd
          3 from sklearn.tree import DecisionTreeClassifier # Import Decision Tree Classifier
          4 from sklearn.model_selection import train_test_split # Import train_test_split function
          5 from sklearn import metrics #Import scikit-learn metrics module for accuracy calculation
```

Loading Data

Let's first load the required Pima Indian Diabetes dataset using pandas' read CSV function. You can download the data [here](#).

```
In [100]: 1 col_names = ['pregnant', 'glucose', 'bp', 'skin', 'insulin', 'bmi', 'pedigree', 'age', 'label']
          2 # Load dataset
          3 pima = pd.read_csv("./data/pima-indians-diabetes.csv", header=None, names=col_names)

In [101]: 1 pima.head()
```

```
Out[101]:
```

	pregnant	glucose	bp	skin	insulin	bmi	pedigree	age	label
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Feature Selection

Here, you need to divide given columns into two types of variables dependent(or target variable) and independent variable(or feature variables).

```
In [102]: 1 #split dataset in features and target variable
2 feature_cols = ['pregnant', 'insulin', 'bmi', 'age', 'glucose', 'bp', 'pedigree']
3 X = pima[feature_cols] # Features
4 y = pima.label # Target variable
```

Splitting Data

To understand model performance, dividing the dataset into a training set and a test set is a good strategy.

Let's split the dataset by using function `train_test_split()`. You need to pass 3 parameters features, target, and test_set size.

```
In [103]: 1 # Split dataset into training set and test set
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1) # 70% training and 30% test/
```

Building Decision Tree Model

Let's create a Decision Tree Model using Scikit-learn.

```
In [104]: 1 # Create Decision Tree classifier object
2 clf = DecisionTreeClassifier()
3
4 # Train Decision Tree Classifier
5 clf = clf.fit(X_train, y_train)
6
7 #Predict the response for test dataset
8 y_pred = clf.predict(X_test)
```

Evaluating Model

Let's estimate, how accurately the classifier or model can predict the type of cultivars.

Accuracy can be computed by comparing actual test set values and predicted values.

```
In [105]: 1 # Model Accuracy, how often is the classifier correct?
2 print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

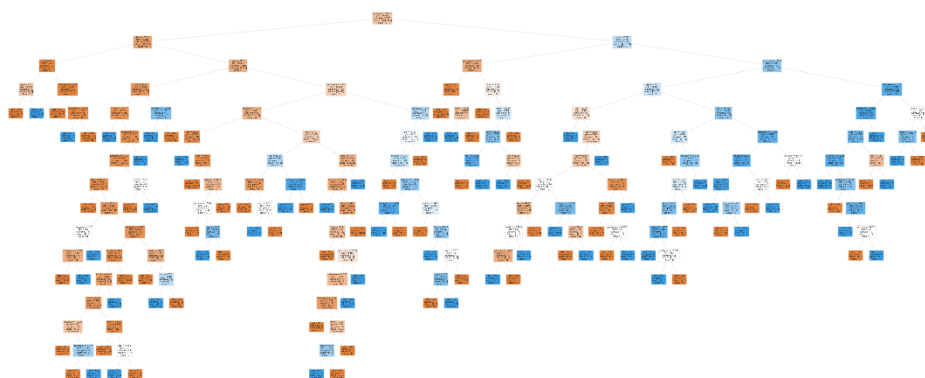
Accuracy: 0.6753246753246753

Well, you got a classification rate of 67.09%, considered as good accuracy. You can improve this accuracy by tuning the parameters in the Decision Tree Algorithm.

Visualizing Decision Trees

You can use Scikit-learn's `plot_tree` function from the tree module to display the tree within a Jupyter notebook.

```
In [107]: 1 from sklearn import tree
2 import matplotlib.pyplot as plt
3
4 plt.figure(figsize=(12, 5), dpi=500)
5 tree.plot_tree(clf,
6               feature_names=feature_cols,
7               class_names=["0", "1"],
8               filled=True,
9               rounded=True);
10 plt.savefig('graphics/dtree.png')
```



Open up `dtree.png` that was generated inside the graphics folder and zoom in to inspect the decision tree chart.

In the decision tree chart, each internal node has a decision rule that splits the data. Gini referred as Gini ratio, which measures the impurity of the node. You can say a node is pure when all of its records belong to the same class, such nodes known as the leaf node.

Here, the resultant tree is unpruned. This unpruned tree is unexplainable and not easy to understand. In the next section, let's optimize it by pruning.

Optimizing Decision Tree Performance

- **criterion** : optional (default="gini") or Choose attribute selection measure: This parameter allows us to use the different-different attribute selection measure. Supported criteria are "gini" for the Gini index and "entropy" for the information gain.
- **splitter** : string, optional (default="best") or Split Strategy: This parameter allows us to choose the split strategy. Supported strategies are "best" to choose the best split and "random" to choose the best random split.
- **max_depth** : int or None, optional (default=None) or Maximum Depth of a Tree: The maximum depth of the tree. If None, then nodes are expanded until all the leaves contain less than `min_samples_split` samples. The higher value of maximum depth causes overfitting, and a lower value causes underfitting ([Source](#)).

In Scikit-learn, optimization of decision tree classifier performed by only pre-pruning. Maximum depth of the tree can be used as a control variable for pre-pruning. In the following the example, you can plot a decision tree on the same data with `max_depth=3`. Other than pre-pruning parameters, You can also try other attribute selection measure such as entropy.

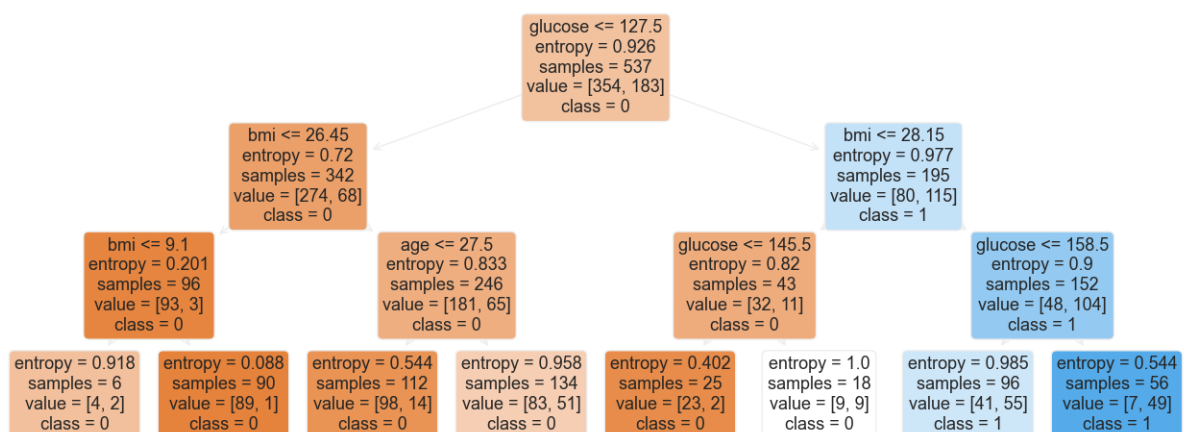
```
In [108]: 1 # Create Decision Tree classifier object
2 clf = DecisionTreeClassifier(criterion="entropy", max_depth=3)
3
4 # Train Decision Tree Classifier
5 clf = clf.fit(X_train,y_train)
6
7 #Predict the response for test dataset
8 y_pred = clf.predict(X_test)
9
10 # Model Accuracy, how often is the classifier correct?
11 print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7705627705627706

Well, the classification rate increased to 77.05%, which is better accuracy than the previous model.

Visualizing Decision Trees again

```
In [109]: 1 plt.figure(figsize=(8, 3), dpi=200)
2 tree.plot_tree(clf,
3               feature_names=feature_cols,
4               class_names=["0", "1"],
5               filled=True,
6               rounded=True);
7 plt.savefig('graphics/diabetes.png')
```



6.4. Model Performance

To measure the performance of a model, we need several elements :


This part is essential

- **Confusion matrix** : also known as the error matrix, allows visualization of the performance of an algorithm :
 - true positive (TP) : Diabetic correctly identified as diabetic
 - true negative (TN) : Healthy correctly identified as healthy
 - false positive (FP) : Healthy incorrectly identified as diabetic
 - false negative (FN) : Diabetic incorrectly identified as healthy
- ***Metrics ***:
 - Accuracy : $(TP + TN) / (TP + TN + FP + FN)$
 - Precision : $TP / (TP + FP)$
 - Recall : $TP / (TP + FN)$
 - F1 score : $2 \times ((Precision \times Recall) / (Precision + Recall))$
- **Roc Curve** : The ROC curve is created by plotting the true positive rate (TPR) against the false positive rate (FPR) at various threshold settings.
- **Precision Recall Curve** : shows the tradeoff between precision and recall for different threshold

Confusion Matrix

```
In [118]: 1 from sklearn.metrics import confusion_matrix
          2 matrix= confusion_matrix(y_test, y_pred)
          3 sns.heatmap(matrix,annot = True, fmt = "d")
```

Out[118]: <AxesSubplot:>



Metrics

Precision

```
In [126]: 1 from sklearn.metrics import precision_score
          2 precision = precision_score(y_test, y_pred)
          3 print("Precision: ",precision)
```

Precision: 0.7105263157894737

recall

```
In [127]: 1 from sklearn.metrics import recall_score
          2 recall = recall_score(y_test, y_pred)
          3 print("Recall is: ",recall)
```

Recall is: 0.6352941176470588

fscore

```
In [128]: 1 print((2*precision*recall)/(precision+recall))

0.6708074534161491
```

This pruned model is less complex, explainable, and easy to understand than the previous decision tree model plot.

Pros

- Decision trees are easy to interpret and visualize.
- It can easily capture Non-linear patterns.
- It requires fewer data preprocessing from the user, for example, there is no need to normalize columns.
- It can be used for feature engineering such as predicting missing values, suitable for variable selection.
- The decision tree has no assumptions about distribution because of the non-parametric nature of the algorithm. ([Source](#))

Cons

- Sensitive to noisy data. It can overfit noisy data.
- The small variation(or variance) in data can result in the different decision tree. This can be reduced by bagging and boosting algorithms.
- Decision trees are biased with imbalance dataset, so it is recommended that balance out the dataset before creating the decision tree.

Conclusion

Congratulations, you have made it to the end of this tutorial!

In this Notebook, I had covered a lot of details about how can do a better EDA and then using a Decision Tree; It's working, attribute selection measures such as Information Gain, Gain Ratio, and Gini Index, decision tree model building, visualization and evaluation on diabetes dataset using Python Scikit-learn package. Also, discussed its pros, cons, and optimizing Decision Tree performance using parameter tuning.

Hopefully, you can now utilize the Decision tree algorithm to analyze your own datasets.

7. Credits

Credits :

- <https://medium.com/@pushkarmandot/https-medium-com-pushkarmandot-what-is-lightgbm-how-to-implement-it-how-to-fine-tune-the-parameters-60347819b7fc>
- https://en.wikipedia.org/wiki/Body_mass_index
- http://rasbt.github.io/mlxtend/user_guide/classifier/EnsembleVoteClassifier/
- <https://www.news-medical.net/health/What-is-Diabetes.aspx>
- <https://lightgbm.readthedocs.io/en/latest/Parameters-Tuning.html>
- <http://ogrisel.github.io/scikit-learn.org/sklearn-tutorial/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>
- <https://www.scikit-yb.org/en/latest/api/classifier/threshold.html>
- <https://www.analyticsindiamag.com/why-is-random-search-better-than-grid-search-for-machine-learning/>
- <https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.LabelEncoder.html>
- https://scikit-learn.org/stable/modules/preprocessing_targets.html#preprocessing-targets
- <https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.StandardScaler.html>
- <https://twitter.com/bearda24>
- <https://www.slideshare.net/DhianaDevaRocha/qcon-rio-machine-learning-for-everyone>
- <https://medium.com/@sebastianorena/some-model-tuning-methods-bfef3e6544f0>
- <https://www.niddk.nih.gov/health-information/diabetes/overview/what-is-diabetes>
- https://en.wikipedia.org/wiki/Pima_people

Thank you all :)